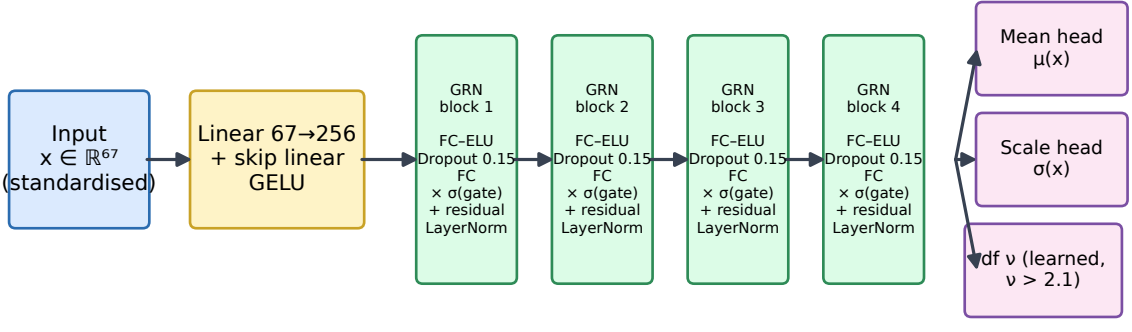

Dropout stays active at inference (Monte Carlo dropout, 60 passes per member); 5 members trained on publication-disjoint folds.
Training loss: Student-t negative log-likelihood + 0.25 × Huber (δ = 0.5) on the mean. No attention mechanism is used.